



Awesome Inc: Universal Sales & Returns

Vaidehi Chaudhari (vdc2009)

Amanpreet Singh Saimbhi (as15798)

Course: ECE-GY-9941 Advance Project III

Project Report

Date of Submission: 5/5/2023

TABLE OF CONTENT

ABSTRACT	3
BUSINESS CASE.....	4
PROJECT MILESTONE	5
Purpose	5
System Requirements.....	5
Roadmap.....	6
ASSUMPTIONS.....	7
OLTP DESIGNS	8
OLTP Logical model.....	8
OLTP Relational Model	8
DATA WAREHOUSING	9
De-normalization.....	9
DW Logical Model.....	9
DW Relational Model	9
ETL SUMMARY.....	10
SQL & DATA ANALYTICS FOR DATA WAREHOUSE	11
ETL Example with CSV file	11
RECORD COUNT OF OLTP TABLE	16
DATA MANAGEMENT	19
External Table	19
History table	19
Table partitioning.....	20
INCREMENTAL ETL	21
DATA VISUALIZATION	23
PROJECT INSIGHTS.....	27
APPENDIX	28
OLTP Data Definition Language Code.....	28
OLTP Data Manipulation Language Code.....	36
OLTP Data dictionary	39
OLAP Data Definition Language Code	41
Procedure code for Incremental ETL	46
History table trigger	48
Partitioning of table.....	48
ETL code	49

ABSTRACT

The goal of this project was to analyze and visualize data using Tableau for effective data-driven decision making. The project involved various tasks, including data extraction, transformation, and loading (ETL), data modeling, and interactive dashboard creation. The data was sourced from multiple data sources and loaded into Snowflake, a cloud-based data warehouse. ETL processes were implemented to clean, transform, and load the data into a suitable format for analysis. Data modeling was done in Snowflake to establish relationships between the data tables and create a foundation for analysis. Finally, interactive dashboards were created using Tableau to visualize the data and provide actionable insights. The dashboards included various charts, tables, and filters, allowing users to explore the data and gain insights into key performance indicators (KPIs). The project successfully achieved the objective of creating an interactive dashboard for data analysis and visualization, which can be used for decision making in the organization.

Keywords: ETL process, Tableau, Snowflake, SQL, data analysis.

BUSINESS CASE

The Universal Sales and Return Data Warehouse and Analytics project aims to centralize sales and return data from multiple sources into a single repository through an ETL process. Currently, Awesome Inc has multiple data sources that contain sales and return data, making it difficult to access and analyze the data effectively. The data is often inconsistent, of low quality, and not easily accessible, leading to inaccurate decision-making and missed business opportunities. This situation creates inefficiencies and can negatively impact the company's bottom line. The proposed solution is to centralize the sales and return data into a single data warehouse using ETL tools such as Oracle Data Modeler. The data will be transformed to ensure consistency, accuracy, and quality before being loaded into the MySQL database.

The transformed data will be analyzed using Tableau for interactive visualizations and dashboards, providing support for data-driven decision-making. The project aims to deliver a comprehensive data warehousing and analytics solution that will enable Awesome Inc to make informed decisions based on accurate and up-to-date data. In addition, the project involves the use of Snowflake, a cloud-based data warehousing platform, for storing and managing large volumes of data. The ETL process involves the extraction of data from different sources, including flat files, databases, and APIs, followed by data cleaning, transformation, and loading into the Snowflake data warehouse. Incremental loading is used to update the data warehouse with new data on a regular basis, ensuring that the data is always current.

The transformed data will be used to create a range of interactive dashboards and visualizations using Tableau, allowing Awesome Inc to gain insights into sales and return trends, customer behavior, and product performance. The project will also enable the creation of ad-hoc reports, enabling quick and easy access to data for business users across the organization.

In short, the Universal Sales and Return Data Warehouse and Analytics project aims to deliver a comprehensive solution that will provide Awesome Inc with a clear view of their sales and return data, enabling informed decision-making and improved business outcomes.

PROJECT MILESTONE

The purpose of this project is to develop a data warehousing solution to extract, transform, and load data into a centralized repository. The goal is to make the data accessible, reliable, and usable for data analytics and reporting. After the ETL process is completed, the project will focus on using Tableau to perform data analytics and gain valuable insights into the data. The outcome of the project is expected to be an enhanced decision-making process by providing a comprehensive understanding of the data.

System Requirements:

Software Requirements:

- Database language: MySQL, Snowflake SQL

Software/Technologies used: Snowflake 3.20.4 , Oracle Data modeler 21.4.2, MySQL Workbench 8.0 CE, AWS S3 bucket, Visual Studio Code, Tableau Desktop 2023.1

Hardware Requirements:

- Windows 10 8GB RAM quad-core processor, Intel i5 processor
- MAC OS X 10.9

Roadmap:



Project Proposal: (23th Jan - 4th Feb)

The project proposal was submitted along with the purpose, objective, system design.



OLTP Design: (5th Feb - 18th Feb)

Logical Database model, Relational Database model, DDL Generation was submitted.



OLTP Database Implementation: (19th Feb - 3rd Mar)

The database was created in MySQL workbench and tables were populated after loading the staging table from CSV files. Also, history table was created to store the deleted entries.



OLAP Design: (4th Mar- 18th March)

Counts of each records were submitted. De-normalized design was created in star schema. Logical model and Relational model of OLAP was submitted.



Data Warehouse design and implementation: (19th Mar - 15th April)

De- normalization from the OLTP design. Use of AWS bucket for S3 bucket to load CSV

Use of Snowflake software which is a cloud computing-based data warehousing software. Made use of Snowflake for cloud-based storage and data analytics. Created history table for address.



ETL code using Snowflake and AWS: (16th Apr – 25th Apr)

The database was created in MySQL workbench and tables were populated after loading the staging table from CSV files. Also, history table was created to store the deleted entries. Created External table in Snowflake. Incremental ETL was also implemented to make sure that the new data is also updated in the data-warehouse.



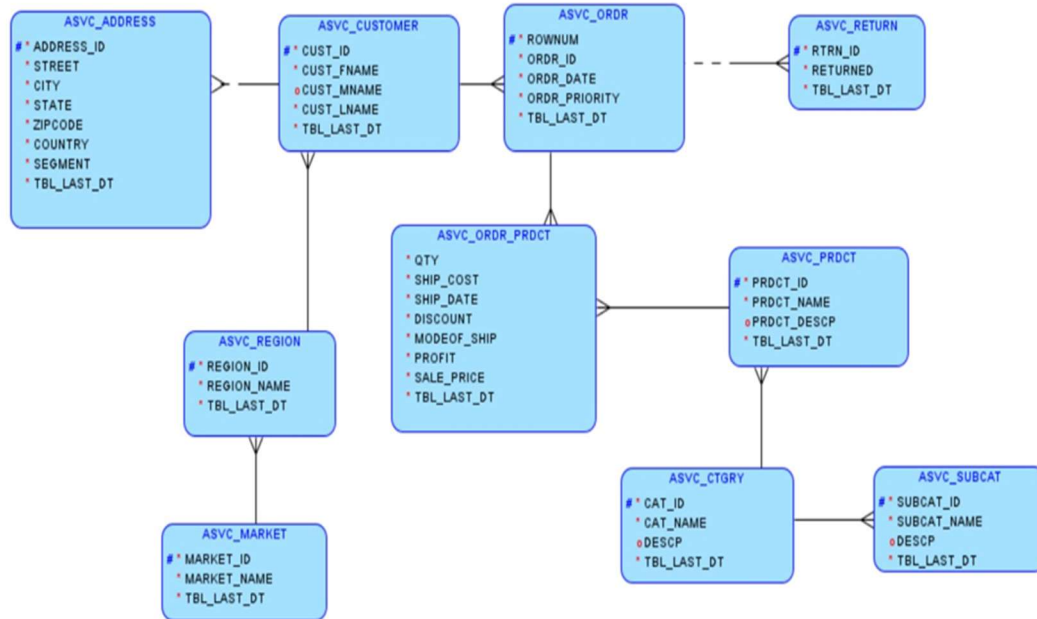
Reports and Analytics from DW database: (26th Apr - 2^d Mar)

Querying of the data and generation of reports and data analytics and data visualization is done by creating dashboards on Tableau Desktop.

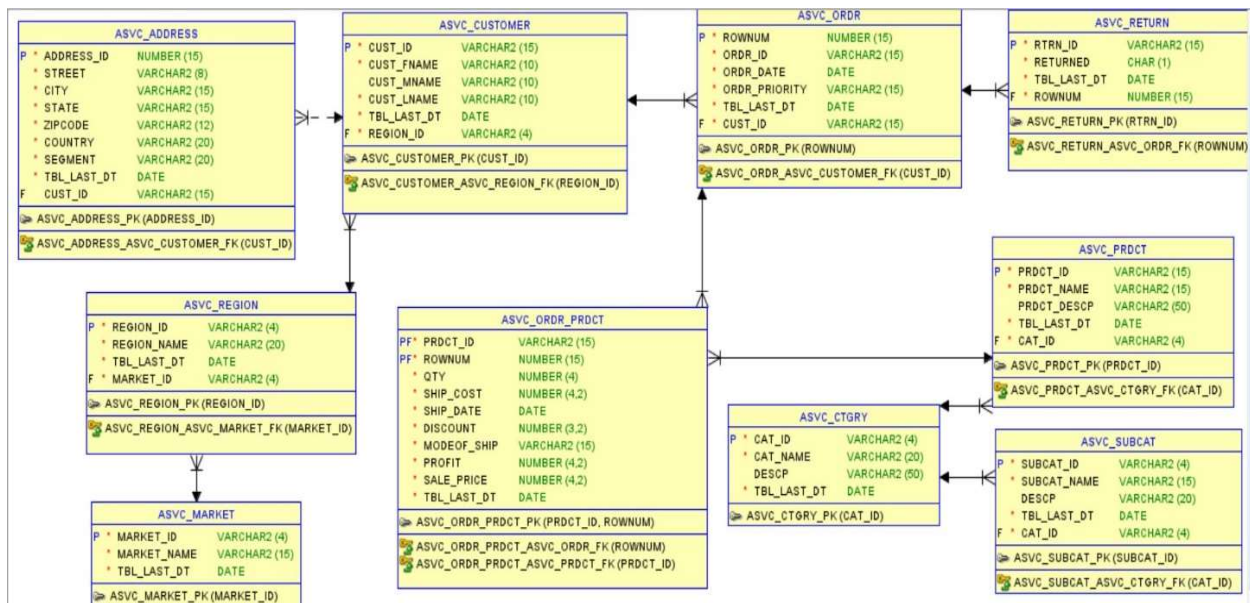
ASSUMPTIONS

1. A customer can have multiple address, but an address belongs to only one customer.
2. A region can have multiple customers, but a customer belongs to only one region.
3. A market can have many regions associated with it but a region will be associated to only one market.
4. Each order is associated with one and only one customer.
5. A customer can place multiple orders.
6. Each order can have multiple products.
7. Each product can be part of multiple orders.
8. Each product belongs to one and only one category.
9. Each category can have multiple products.
10. Each category can have multiple sub-categories.
11. Each sub-category belongs to one and only one category.
12. The region and market details of a customer are based on the shipping address of their orders.
13. The sales price, shipping cost, and profit details are associated with each individual product in an order.

OLTP LOGICAL MODEL



OLTP RELATIONAL MODEL

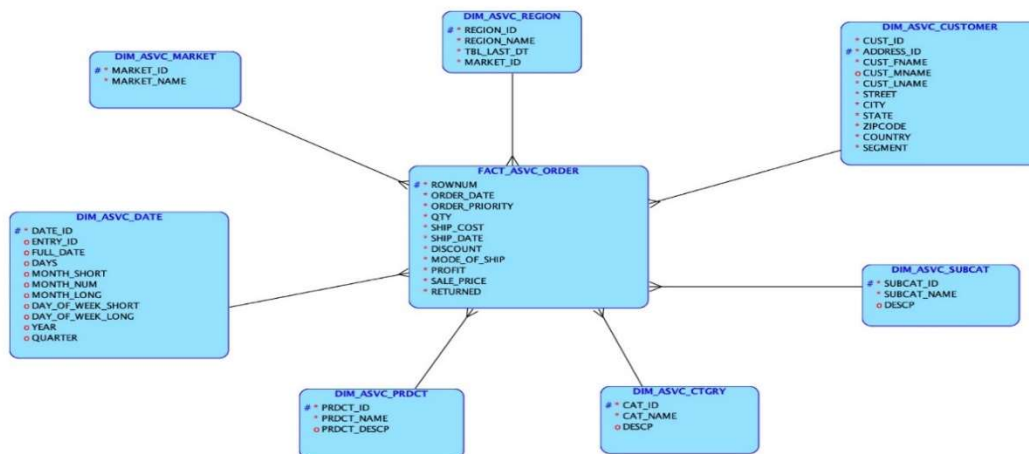


DATA WAREHOUSING

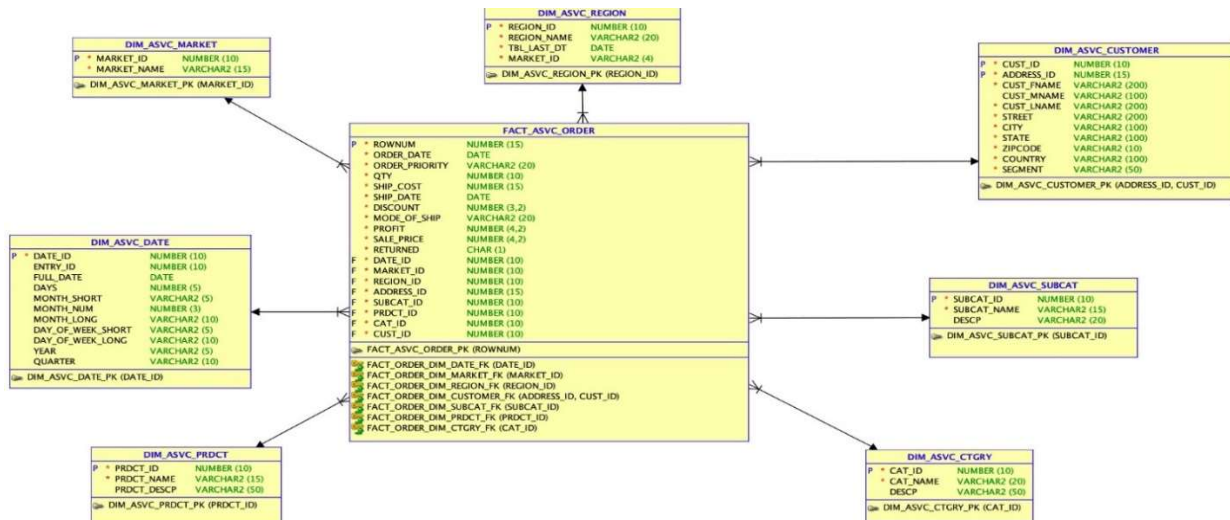
Data warehousing De-Normalizing:

Most of the modern applications need to be able to retrieve data in the shortest time possible, so that's why we de-normalize the tables to have fewer join. However, updates and inserts are more expensive, denormalization can make update and insert code harder to read, and it also might update and insert code harder to write.

OLAP LOGICAL MODEL



OLAP RELATIONAL MODEL



ETL SUMMARY

The ETL process is a key part of the project as it enables us to extract data from the source system (an S3 bucket), transform it into a more usable format, and load it into our target system (Snowflake). The ETL process typically consists of three phases: Extract, Transform, and Load. Here's how these phases of ETL to our project:

Extract:

In this phase, we extract data from the S3 bucket and make it available for further processing. We achieved this by creating an external stage in Snowflake that is linked to the S3 bucket via an AWS S3 Integration. This stage serves as the interface between our source data and our target system (Snowflake).

Transform:

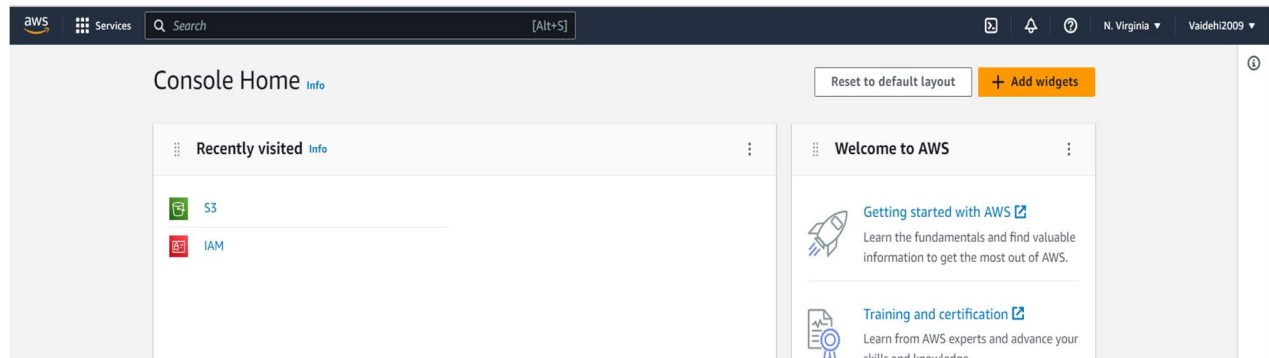
In this phase, we transform the data into a more usable format. This often involves filtering, aggregating, and joining data from different sources. In our project, we need to transform the data to conform to a specific schema or format that is required by the business needs.

Load:

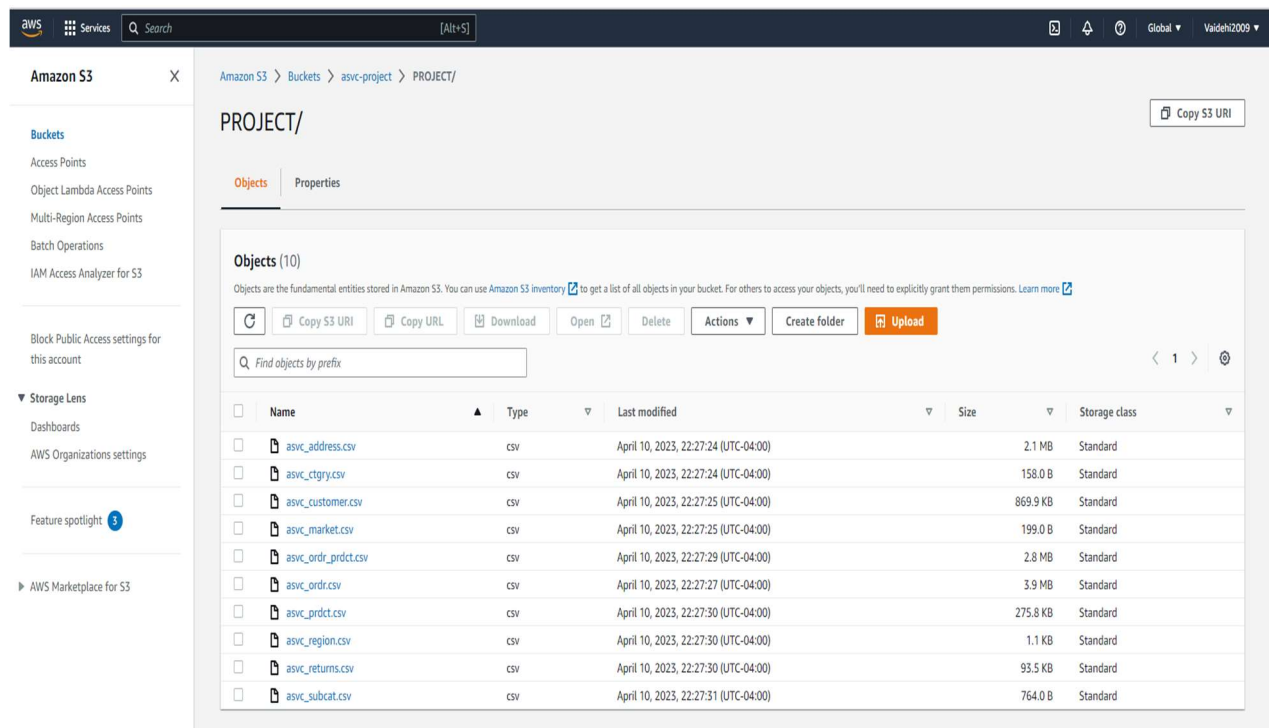
In this phase, we load the transformed data into our target system (Snowflake). We accomplished this by creating a Snowflake stage that is linked to the external stage in the previous step. This stage serves as the target location in Snowflake where data will be loaded from the external S3 stage during the ETL process. Once the data is loaded, we will use it for reporting, analysis, and other business needs.

SQL AND DATA ANALYTICS FOR DATA WAREHOUSE

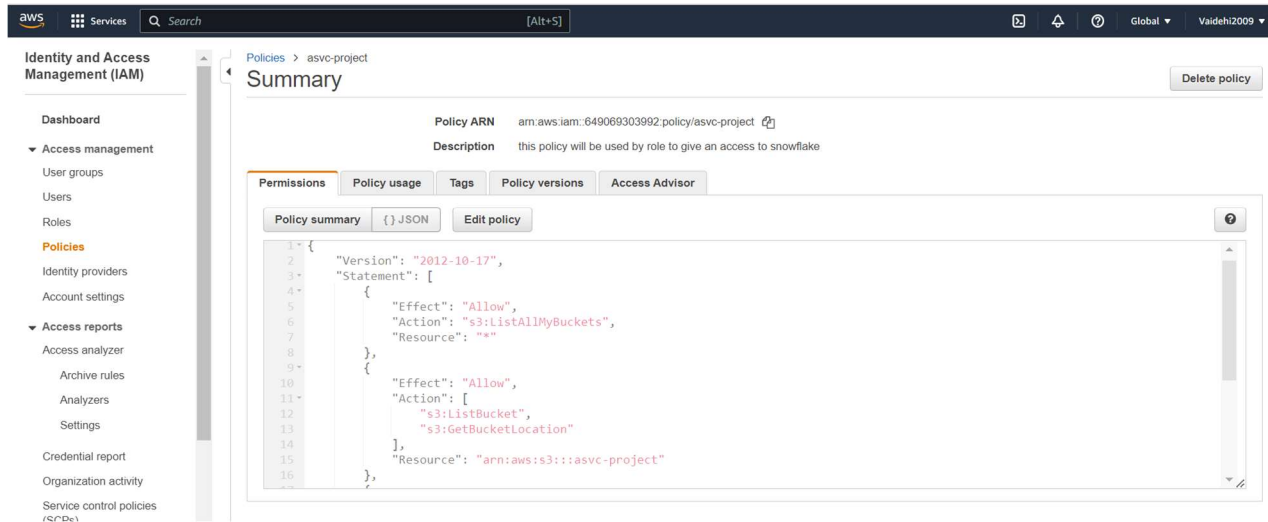
Created AWS S3 bucket and Wrote a JSON object to define an IAM policy that allows ETL tools to read from source S3 bucket, transform the data and write the results to S3 bucket.



Uploaded the csv file of OLTP to the AWS bucket:



Create connection between AWS and snowflake, so snowflake can access the data in AWS bucket.



Generated de-normalized OLAP for data warehouse in which ETL process will be performed and the data from OLTP will be loaded.



This is a key step in the ETL process, specifically the "Extract" phase, where data is extracted from the source system (in this case, an S3 bucket) and made available for further processing.

Created a storage integration, this sets up the integration between Snowflake and the S3 bucket, allowing Snowflake to access and load data from the external stage.

The screenshot shows the Snowflake SQL Editor interface. On the left, the 'Databases' tab is active, showing a tree view with 'ASVC_PROJECT' expanded. The main editor area displays SQL code for creating a schema and an external stage integration. A red box highlights the 'CREATE OR REPLACE STORAGE integration' block.

```
1 create schema asvc_project
2
3
4 CREATE OR REPLACE STORAGE integration aws_s3_integration
5 type = external_stage
6 storage_provider='S3'
7 enabled=true
8 storage_aws_role_arn= 'arn:aws:iam::649069303992:role/asvc-project'
9 storage_allowed_locations= ('s3://asvc-project/');
10
11 SHOW INTEGRATIONS;
12
13 DESC integration aws_s3_integration;
14
15 GRANT USAGE ON INTEGRATION aws_s3_integration TO ROLE accountadmin;
```

Created a Snowflake Stage which sets up a target location in Snowflake where data will be loaded from the external S3 stage during the ETL process:

The screenshot shows the Snowflake SQL Editor interface. The main editor area displays SQL code for creating a stage. The code is as follows:

```
19 field_delimiter='|'
20 skip_header=1;
21
22 CREATE OR REPLACE STAGE demo_aws_stage
23 storage_integration = aws_s3_integration
24 file_format = demo_format
25 url='s3://asvc-project/';
26
27 List @demo_aws_stage
28
```

Loading of the unstructured data from CSV file to the OLAP table dim_asvc_category:

The screenshot shows the Snowflake SQL Editor interface. The main editor area displays SQL code for loading data from a CSV file into a table. The code is as follows:

```
33 );
34
35
36 COPY INTO dim_asvc_ctgry
37 FROM @demo_aws_stage/PROJECT/asvc_ctgry.csv
38 FILE_FORMAT=(format_name=demo_format1 error_on_column_count_mismatch=false)
39 ON_ERROR = 'CONTINUE';
40
41 select * from dim_asvc_ctgry;
```

Below the editor, the 'Results' tab is active, showing a table with 3 rows and 3 columns: CAT_ID, CAT_NAME, and DESC. The table contains the following data:

CAT_ID	CAT_NAME	DESC
1	"M001"	"Furniture"
2	"M002"	"Office Supplies"
3	"M003"	"Technology"

On the right side of the results, the 'Query Details' panel shows the query duration as 1.1s and the number of rows as 3.

Loading of the unstructured data from CSV file to the OLAP table dim_asvc_customer:

The screenshot shows a data visualization tool interface. At the top, there's a header with 'Worksheets', a timestamp '2023-04-10 12:30pm', and a 'PREVIEW' button. Below the header, there's a sidebar on the left with 'Databases' and 'Worksheets' tabs. The 'Databases' tab is active, showing a tree view of databases: 'ASVC_PROJECT', 'SNOWFLAKE', and 'SNOWFLAKE_SAMPLE_DATA'. The main area displays a query editor with the SQL query: `select * from dim_asvc_customer`. Below the query editor, there's a 'Results' tab showing a table with 11 rows and 10 columns: CUST_ID, ADDRESS_ID, CUST_FNAME, CUST_MNAME, CUST_LNAME, STREET, CITY, STATE, ZIPCODE, and COUNTRY. The table contains data for various customers, including Aaron Bergman, Justin Ritter, Craig Reiter, Katherine Murray, Rick Hansen, Jim Mitchum, Toby Swindell, Mick Brown, Aaron Bergman, Anthony Jacobs, and Magdelene Morse. To the right of the table, there's a 'Query Details' panel showing 'Query duration: 1.2s' and 'Rows: 10K'. Below this, there are two bar charts: one for 'CUST_ID' showing '100% filled' and another for 'ADDRESS_ID' showing a distribution from 16 to 11467. The bottom of the interface has a 'Results' tab and a 'Chart' tab.

	CUST_ID	ADDRESS_ID	CUST_FNAME	CUST_MNAME	CUST_LNAME	STREET	CITY	STATE	ZIPCODE	COUNTRY
1	AB-100151402	16	Aaron		Bergman		Oklahoma City	Oklahoma	73120	United Stat
2	JR-162107	17	Justin		Ritter		Wollongong	New South Wales		Australia
3	CR-127307	18	Craig		Reiter		Brisbane	Queensland		Australia
4	KM-1637548	19	Katherine		Murray		Berlin	Berlin		Germany
5	RH-9495111	20	Rick		Hansen		Dakar	Dakar		Senegal
6	JM-156557	21	Jim		Mitchum		Sydney	New South Wales		Australia
7	TS-2134092	22	Toby		Swindell		Porirua	Wellington		New Zealar
8	MB-1808592	23	Mick		Brown		Hamilton	Waikato		New Zealar
9	AB-100151404	24	Aaron		Bergman		Seattle	Washington	98103	United Stat
10	AJ-107801	25	Anthony		Jacobs		Kabul	Kabul		Afghanistan
11	MM-7260110	26	Magdelene		Morse		Jizan	Jizan		Saudi Arabi

Retrieved information about all tables in the specified schema that have been altered within the last 7 days, sorted in descending order based on the date they were last altered.

The screenshot shows a data visualization tool interface. At the top, there's a header with 'Worksheets', a timestamp '2023-04-10 12:30pm', and a 'PREVIEW' button. Below the header, there's a sidebar on the left with 'Databases' and 'Worksheets' tabs. The 'Databases' tab is active, showing a tree view of databases: 'ASVC_PROJECT', 'SNOWFLAKE', and 'SNOWFLAKE_SAMPLE_DATA'. The main area displays a query editor with the SQL query: `SELECT * FROM INFORMATION_SCHEMA.TABLES WHERE TABLE_SCHEMA = 'ASVC_PROJECT' AND LAST_ALTERED >= DATEADD(day, -7, CURRENT_TIMESTAMP()) ORDER BY LAST_ALTERED DESC;`. Below the query editor, there's a 'Results' tab showing a table with 8 rows and 7 columns: TABLE_CATALOG, TABLE_SCHEMA, TABLE_NAME, TABLE_OWNER, TABLE_TYPE, and IS_TRANSIENT. The table contains information about tables in the ASVC_PROJECT schema, including ASVC_ORDR_PRDCT, FACT_ASVC_ORDER, DIM_ASVC_SUBCAT, DIM_ASVC_REGION, DIM_ASVC_PRDCT, DIM_ASVC_MARKET, DIM_ASVC_DATE, and DIM_ASVC_CUSTOMER. To the right of the table, there's a 'Query Details' panel showing 'Query duration: 1.2s' and 'Rows: 10K'. Below this, there are two bar charts: one for 'CUST_ID' showing '100% filled' and another for 'ADDRESS_ID' showing a distribution from 16 to 11467. The bottom of the interface has a 'Results' tab and a 'Chart' tab.

	TABLE_CATALOG	TABLE_SCHEMA	TABLE_NAME	TABLE_OWNER	TABLE_TYPE	IS_TRANSIENT
1	ASVC_PROJECT	ASVC_PROJECT	ASVC_ORDR_PRDCT	ACCOUNTADMIN	EXTERNAL TABLE	null
2	ASVC_PROJECT	ASVC_PROJECT	FACT_ASVC_ORDER	ACCOUNTADMIN	BASE TABLE	NO
3	ASVC_PROJECT	ASVC_PROJECT	DIM_ASVC_SUBCAT	ACCOUNTADMIN	BASE TABLE	NO
4	ASVC_PROJECT	ASVC_PROJECT	DIM_ASVC_REGION	ACCOUNTADMIN	BASE TABLE	NO
5	ASVC_PROJECT	ASVC_PROJECT	DIM_ASVC_PRDCT	ACCOUNTADMIN	BASE TABLE	NO
6	ASVC_PROJECT	ASVC_PROJECT	DIM_ASVC_MARKET	ACCOUNTADMIN	BASE TABLE	NO
7	ASVC_PROJECT	ASVC_PROJECT	DIM_ASVC_DATE	ACCOUNTADMIN	BASE TABLE	NO
8	ASVC_PROJECT	ASVC_PROJECT	DIM_ASVC_CUSTOMER	ACCOUNTADMIN	BASE TABLE	NO

Continuation

```

99
100 SELECT *
101 FROM INFORMATION_SCHEMA.TABLES
102 WHERE TABLE_SCHEMA = 'ASVC_PROJECT'
103 AND LAST_ALTERED >= DATEADD(day, -7, CURRENT_TIMESTAMP())
104 ORDER BY LAST_ALTERED DESC;
105
106
107
108
109
110

```

Results
Chart

		CREATED	LAST_ALTERED	LAST_DDL	LAST_DDL_BY	...	AUTO_CLUST
1		2023-04-14 14:21:33.588 -0700	2023-04-14 14:21:34.124 -0700	2023-04-14 14:21:33.588 -0700	VDC2009		NO
2		2023-04-11 19:27:56.764 -0700	2023-04-11 19:28:01.134 -0700	2023-04-11 19:28:01.134 -0700	VDC2009		NO
3		2023-04-11 19:27:44.582 -0700	2023-04-11 19:27:48.204 -0700	2023-04-11 19:27:48.204 -0700	VDC2009		NO
4		2023-04-11 19:27:36.965 -0700	2023-04-11 19:27:40.460 -0700	2023-04-11 19:27:40.460 -0700	VDC2009		NO
5		2023-04-11 19:27:30.014 -0700	2023-04-11 19:27:33.157 -0700	2023-04-11 19:27:33.157 -0700	VDC2009		NO
6		2023-04-11 19:27:21.685 -0700	2023-04-11 19:27:24.760 -0700	2023-04-11 19:27:24.760 -0700	VDC2009		NO
7		2023-04-11 19:07:49.358 -0700	2023-04-11 19:07:52.987 -0700	2023-04-11 19:07:52.987 -0700	VDC2009		NO
8		2023-04-11 11:31:41.242 -0700	2023-04-11 11:32:15.460 -0700	2023-04-11 11:31:41.242 -0700	VDC2009		NO

Query Details

Query duration
1.6s

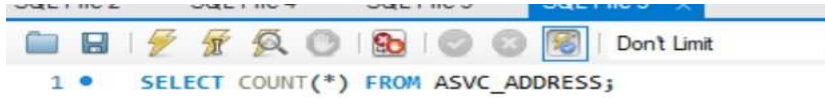
Rows
11

TABLE_CATALOG
ASVC_PROJECT

TABLE_SCHEMA
ASVC_PROJECT

RECORD COUNT OF ALL TABLES IN OLTP

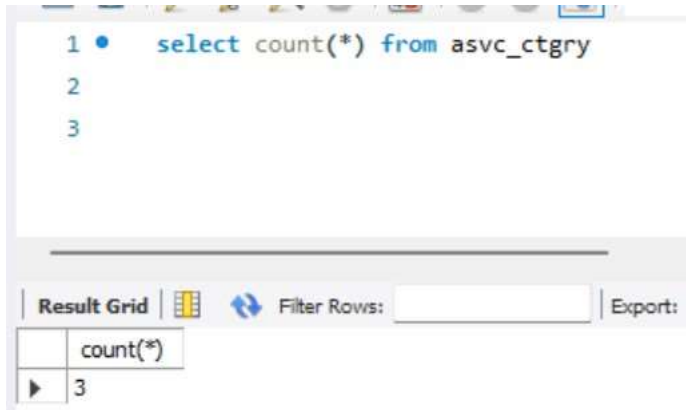
Count for asvc_address:



```
1 • SELECT COUNT(*) FROM ASVC_ADDRESS;
```

COUNT(*)
25178

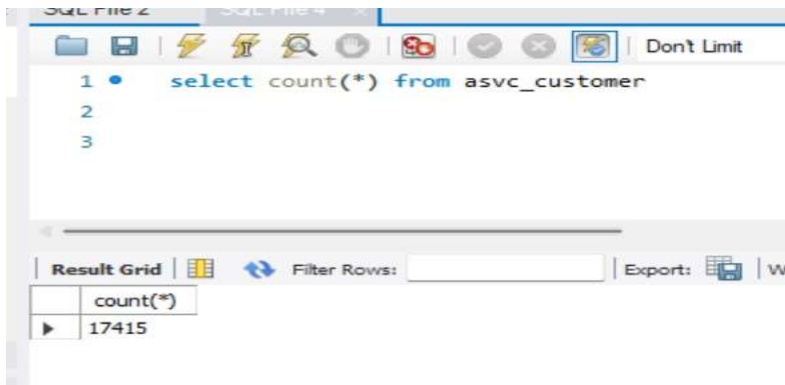
Count for asvc_ctgry:



```
1 • select count(*) from asvc_ctgry
2
3
```

count(*)
3

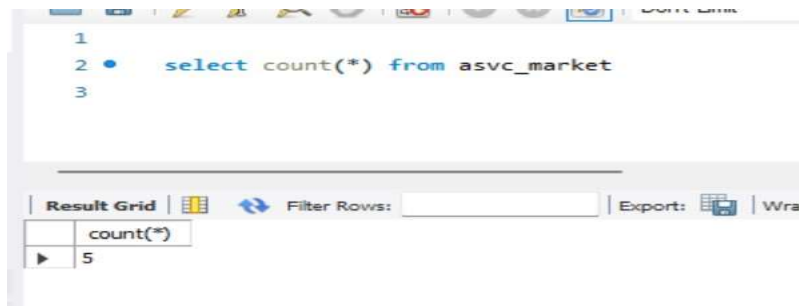
Count for asvc_customer:



```
1 • select count(*) from asvc_customer
2
3
```

count(*)
17415

Count for asvc_market:



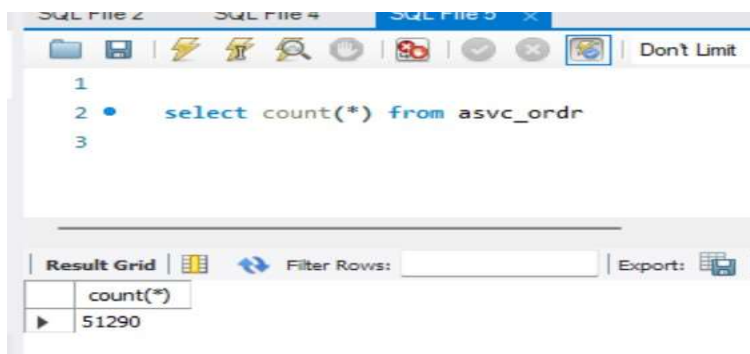
The screenshot shows a SQL IDE window with a query editor containing the following SQL statement:

```
1  
2 • select count(*) from asvc_market  
3
```

Below the query editor, the 'Result Grid' tab is active, displaying the following table:

	count(*)
▶	5

Count for asvc_ordr:



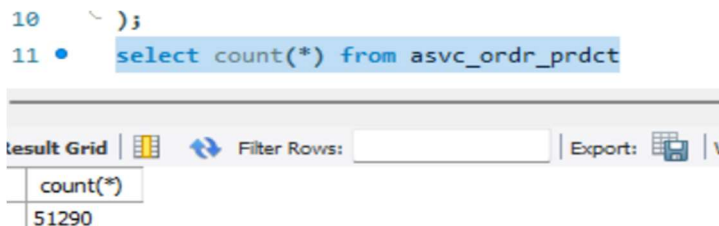
The screenshot shows a SQL IDE window with a query editor containing the following SQL statement:

```
1  
2 • select count(*) from asvc_ordr  
3
```

Below the query editor, the 'Result Grid' tab is active, displaying the following table:

	count(*)
▶	51290

Count for asvc_ordr_prdct:



The screenshot shows a SQL IDE window with a query editor containing the following SQL statement:

```
10 • );  
11 • select count(*) from asvc_ordr_prdct
```

Below the query editor, the 'Result Grid' tab is active, displaying the following table:

	count(*)
▶	51290

Count for asvc_return:



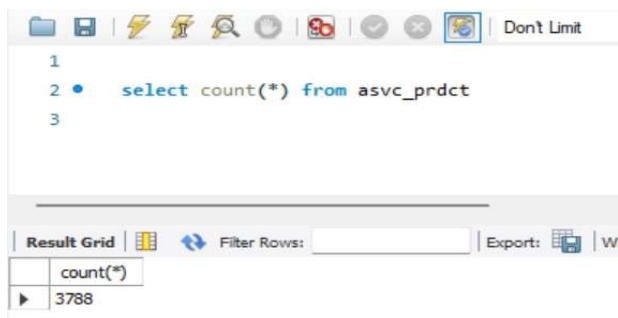
The screenshot shows a SQL IDE window with a query editor containing the following SQL statement:

```
1  
2 • select count(*) from asvc_return;  
3
```

Below the query editor, the 'Result Grid' tab is active, displaying the following table:

	count(*)
▶	2220

Count for asvc_prdct:



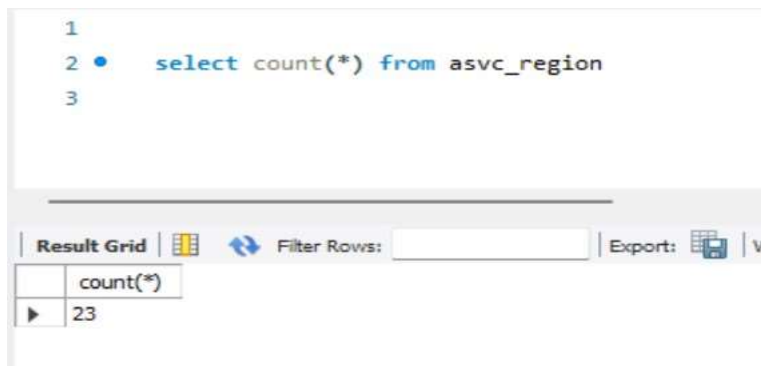
The screenshot shows a SQL query editor with the following text:

```
1  
2 • select count(*) from asvc_prdct  
3
```

Below the editor is a toolbar with icons for file operations, a search icon, and a 'Don't Limit' button. Below the toolbar is a 'Result Grid' section with a 'Filter Rows' input field and an 'Export' button. The result grid contains the following data:

count(*)
3788

Count for asvc_region:




The screenshot shows a SQL query editor with the following text:

```
1  
2 • select count(*) from asvc_region  
3
```

Below the editor is a toolbar with icons for file operations, a search icon, and a 'Don't Limit' button. Below the toolbar is a 'Result Grid' section with a 'Filter Rows' input field and an 'Export' button. The result grid contains the following data:

count(*)
23

Count for asvc_return:



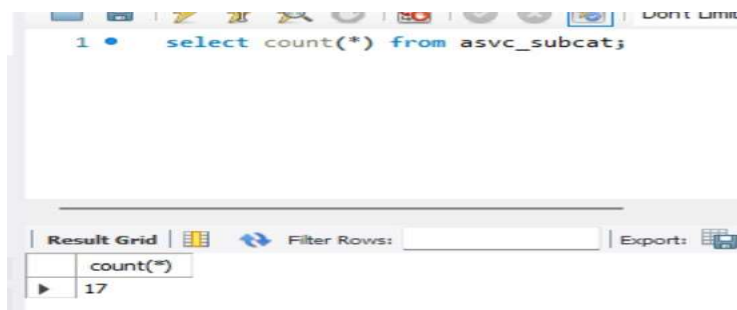
The screenshot shows a SQL query editor with the following text:

```
1  
2 • select count(*) from asvc_return;  
3
```

Below the editor is a toolbar with icons for file operations, a search icon, and a 'Don't Limit' button. Below the toolbar is a 'Result Grid' section with a 'Filter Rows' input field and an 'Export' button. The result grid contains the following data:

count(*)
2220

Count for asvc_subcat:



The screenshot shows a SQL query editor with the following text:

```
1 • select count(*) from asvc_subcat;
```

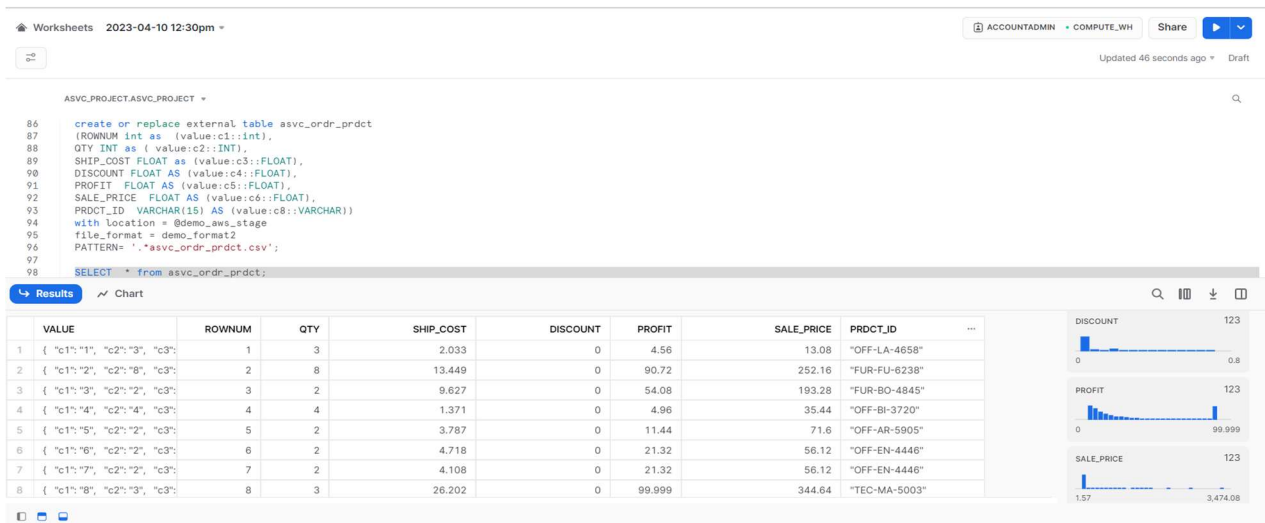
Below the editor is a toolbar with icons for file operations, a search icon, and a 'Don't Limit' button. Below the toolbar is a 'Result Grid' section with a 'Filter Rows' input field and an 'Export' button. The result grid contains the following data:

count(*)
17

DATA MANAGEMENT

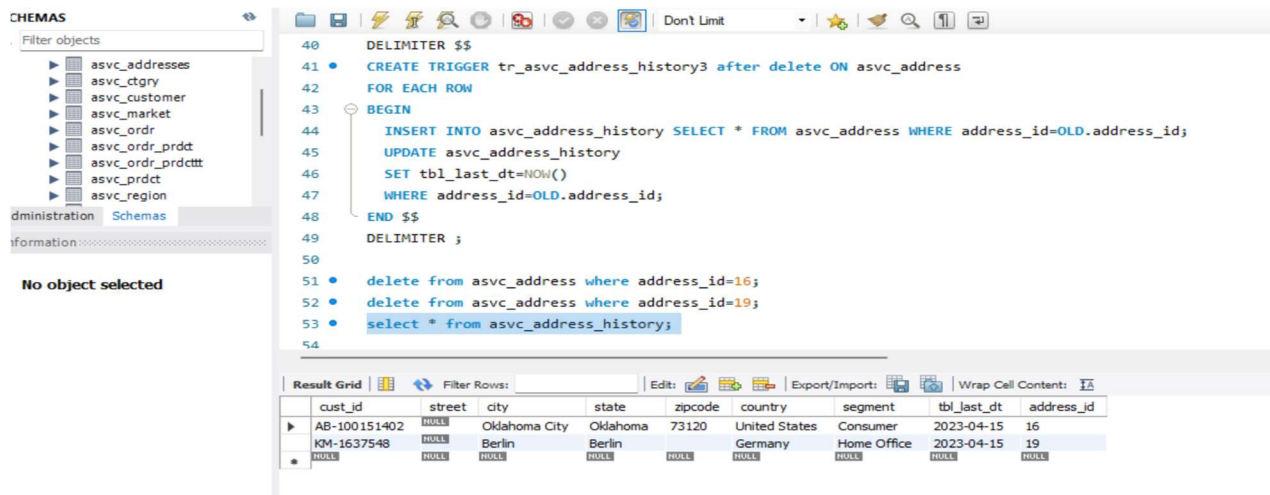
EXTERNAL TABLE:

Created an “**External table**” that query the data in the CSV files stored in S3 as if they were regular Snowflake tables, without the need to load the data into Snowflake first.



HISTORY TABLE:

Created history table in OLTP for asvc_address where all the deleted records will be stored in history table.



PARTITION TABLE

In Snowflake, table partitioning is not supported in the traditional sense, but clustering is a feature that serves a similar purpose. Table clustering is the process of organizing data within a table based on the values in one or more columns. It reorders the data in the table so that rows with similar values are stored together physically, allowing for more efficient data retrieval.

When a table is clustered on one or more columns, Snowflake uses those columns to group together data in the same micro-partitions. By clustering on specific columns, Snowflake can ensure that rows with similar values will be stored in the same micro-partitions, which can make queries more efficient by minimizing the amount of data that needs to be scanned.

We have used clustering which serves the same purpose as table partitioning

```
30
31
32 CREATE TABLE partitioning_fact_asvc cluster by (MARKET_ID) AS
33 (SELECT * FROM fact_asvc_order);
34
35 SELECT SYSTEM$CLUSTERING_INFORMATION('partitioning_fact_asvc');
36
37 SHOW TABLES LIKE 'partitioning_fact_asvc';
38
39
40 -- Code to verify that the number of clusters are dependent on data.
41 -- CREATE OR REPLACE TABLE CLUSTERING_DEMO cluster by (TENANT_ID) (
42 -- ID NUMBER(10,0),
43 -- TENANT_ID NUMBER(4,0),
44
Results Chart
```

SYSTEM\$CLUSTERING_INFORMATION('PARTITIONING_FACT_ASVC')
{ "cluster_by_keys": "LINEAR(MARKET_ID)", "total_partition_count": 1, "total_constant_partition_count": 0, "average_overlaps": 0.0, "average_depth": 1.0,

In our project, we have set the Market ID as the partition ID because we have only five markets, which will essentially create five clusters of data. This will enable us to search for specific data in a particular cluster according to the market. However, Snowflake's internal implementation of partitioning is dependent on the size of our data and available nodes.

Since we have around 50k rows, it is not sufficient for Snowflake to automatically create partitions. Therefore, the partition count is still coming out to be one. However, Snowflake will automatically make partitions in the future if the data scales to a larger extent. To verify this, we have created a dummy table and filled it with random data, and we found that Snowflake created partitions when the data exceeded a certain threshold.

INCREMENTAL ETL

Incremental ETL is a method of data extraction and processing that involves extracting only the data that has changed since the last ETL run, transforms it into the desired format, and loads it into the target system, rather than processing the entire data set each time. This method is often used to reduce processing time and improve efficiency, especially for large and constantly changing data sets.

For demo, we are changing the order_priority from medium to high for a particular order_id in order to update these changes in the Snowflake schema.

```
-- Updating from medium to high.
UPDATE asvc_ordr SET ordr_priority = "High", tbl_last_dt = curdate()
WHERE ordr_id = "AE-2012-P08865138-41184";
```

This query selects records modified in the last 24 hours. The results are written to a CSV file, which is then pushed in AWS S3 bucket.

- ```
select 'ordr_ID', 'ORDR_DATE', 'ORDR_PRIORITY', 'QTY', 'SHIP_DATE', 'SHIP_COST', 'MODE_OF_SHIP', 'DISCOUNT',
'ADDRESS_ID', 'MARKET_ID', 'REGION_ID', 'PRDCT_ID', 'CAT_ID', 'CUST_ID', 'subcat_id', 'subcat_name', 'returned
UNION
SELECT ordr_prd.ordr_ID, ordr.ORDR_DATE, ordr.ORDR_PRIORITY, ordr_prd.QTY, ordr_prd.SHIP_DATE,
ordr_prd.SHIP_COST, ordr_prd.MODE_OF_SHIP, ordr_prd.DISCOUNT, ordr_prd.PROFIT, ordr_prd.SALE_PRICE, adr
mar.MARKET_ID, reg.REGION_ID, prod.PRDUCT_ID, ctg.CAT_ID, cust.CUST_ID, sub.subcat_id, sub.subcat_name,
ret.returned
FROM ASVC_ORDR ordr
JOIN ASVC_ORDR_PRDCT ordr_prd ON ordr.ordr_id = ordr_prd.ORDR_ID
JOIN ASVC_PRDCT prod ON prod.prdct_id = ordr_prd.prdct_id
JOIN ASVC_CUSTOMER cust ON ordr.cust_id = cust.cust_id
JOIN ASVC_CTGRY ctg ON ctg.cat_id = prod.cat_id
JOIN ASVC_SUBCAT sub ON sub.subcat_id = prod.subcat_id
JOIN ASVC_ADDRESS addr ON addr.address_id = ordr.address_id
JOIN ASVC_RETURNS ret ON ordr.ordr_id = ret.ordr_id
JOIN ASVC_REGION reg ON cust.region_id = reg.region_id
JOIN ASVC_MARKET mar ON reg.market_id = mar.market_id
WHERE ordr.tbl_last_dt > date_sub(curdate(), interval 1 day)
INTO OUTFILE 'C:/ProgramData/MySQL/MySQL Server 8.0/Uploads/incremental.csv'
FIELDS TERMINATED BY ',' OPTIONALLY ENCLOSED BY '"' LINES TERMINATED BY '\n';
```

After uploading the CSV file in AWS S3 bucket, we wrote a procedure which will update the new data into the existing fact\_table.

```
CREATE OR REPLACE PROCEDURE load_merge_fact_asvc_order()
RETURNS VARCHAR
LANGUAGE JAVASCRIPT
AS $$
 var err_code;
 var err_msg;
 try {
 var stmt = snowflake.createStatement({sqlText: `
 MERGE INTO fact_asvc_order a
 USING stg_incrmntl_fact_table b
 ON (a.ordr_id = b.ordr_id) AND (a.prdct_id = b.prdct_id)
 WHEN MATCHED THEN
 UPDATE SET
 a.order_date = b.order_date,
 a.order_priority = b.order_priority,
 a.qty = b.qty,
 a.ship_date = b.ship_date,
 a.ship_cost = b.ship_cost,
 a.mode_of_ship = b.mode_of_ship,
 a.discount = b.discount,
 a.profit = b.profit,
 a.sale_price = b.sale_price,
 a.address_id = b.address_id,
 a.market_id = b.market_id,
 a.region_id = b.region_id,
 a.cat_id = b.cat_id,
 a.cust_id = b.cust_id,
 a.subcat_id = b.subcat_id,
 a.subcat_name = b.subcat_name,
 a.returned = b.returned
 WHEN NOT MATCHED THEN
 INSERT (ordr_id, order_date, order_priority, qty, ship_date, ship_cost, mode_of_ship, discount, profit,
 sale_price, address_id, market_id, region_id, prdct_id, cat_id, cust_id, subcat_id, subcat_name, returned)
 VALUES (b.ordr_id, b.order_date, b.order_priority, b.qty, b.ship_date, b.ship_cost, b.mode_of_ship, b.discount,
 b.profit, b.sale_price, b.address_id, b.market_id, b.region_id, b.prdct_id, b.cat_id, b.cust_id, b.subcat_id, b.subcat_name, b.returned);
 `});
 var result = stmt.execute();
 return result.next() ? result.getColumnValue(1).toString() : 'Success';
 } catch (err) {
 err_code = err.code;
 err_msg = err.message;
 return 'Error code ' + err_code + ': ' + err_msg;
 }
$$;

CALL load_merge_fact_asvc_order();
```

Now as we can see in the below screenshot that the order\_priority has changed from medium to high.

```
321 CALL load_merge_fact_asvc_order();
322
323
324 select * from fact_asvc_order
325 where ord_r_id= 'AE-2012-P08865138-41184'
326
327
```

|   | ORDR_ID                 | ORDER_DATE | ORDER_PRIORITY | QTY | SHIP_DATE  | SHIP_COST | MODE_OF_SHIP   | DISCOUNT |
|---|-------------------------|------------|----------------|-----|------------|-----------|----------------|----------|
| 1 | AE-2012-P08865138-41184 | 2012-10-02 | High           | 2   | 2012-10-06 | 5.69      | Standard Class | 0.7      |
| 2 | AE-2012-P08865138-41184 | 2012-10-02 | High           | 6   | 2012-10-06 | 3.87      | Standard Class | 0.7      |

# DATA VISUALIZATION

Data Visualization is a powerful tool that can help you convey complex information in a clear and efficient manner. By using graphics such as charts, tables, and infographics, you can make large amounts of data more accessible and understandable. With the help of visual analytics and dashboards, you can analyze raw data and gain insights into your business operations, both past and future.

By sharing your findings and monitoring progress, you can make better strategic decisions and identify new opportunities. Overall, data visualization can greatly enhance your ability to understand and communicate important information.

We have used Tableau Desktop 2023.1 as a Data visualization tool for creating charts and graphs as it is easy to use and low barrier to entry.

We created a link between Snowflake and Tableau so that Tableau can directly access the data and the schema that is present in the Snowflake data warehouse.

The screenshot displays the Tableau Desktop 2023.1 interface. On the left, the 'Connections' pane is highlighted with a red box, showing a connection to 'slb42968.us-...omputing.com' (Snowflake). Below this, the 'Warehouse' is set to 'COMPUTE\_WH', the 'Database' is 'ASVC\_PROJECT', and the 'Schema' is 'ASVC\_PROJECT'. The 'Schema' list includes 'DIM\_ASVC...USTOMER)', 'DIM ASVC... SVC DATE)', 'New Custom SOI', 'New Union', and 'New Table Extension'.

The main workspace shows a data model diagram with a central box labeled 'FACT\_ASVC\_ORDER' connected to three other boxes: 'DIM\_ASVC\_CUSTOMER', 'DIM\_ASVC\_DATE', and 'DIM\_ASVC\_MARKET'. Below the diagram, a table view shows data for 'FACT\_ASVC\_ORDER' with columns: 'Order Id', 'Order Date', 'Order Priority', and 'QTY'. The table contains three rows of data.


| Order Id                 | Order Date | Order Priority | QTY |
|--------------------------|------------|----------------|-----|
| MX-2014-BE1145531-41650  | 1/11/2014  | Medium         | 6   |
| US-2012-JF15490140-41269 | 12/26/2012 | Medium         | 3   |
| ES-2015-CS1235548-42010  | 1/6/2015   | Medium         | 2   |

The bottom of the screen shows the Windows taskbar with various application icons and the system clock indicating 4:20 PM on 4/26/2023.



### Worksheet 1: Sales growth over months for the category.

Sales growth over months for the category Furniture, office supplies, and technology. This graph is necessary because we can quickly see how sales are changing over time, which categories are performing better or worse, and where there might be opportunities to improve performance.



### Worksheet 2: Sub-Category Profit Above average

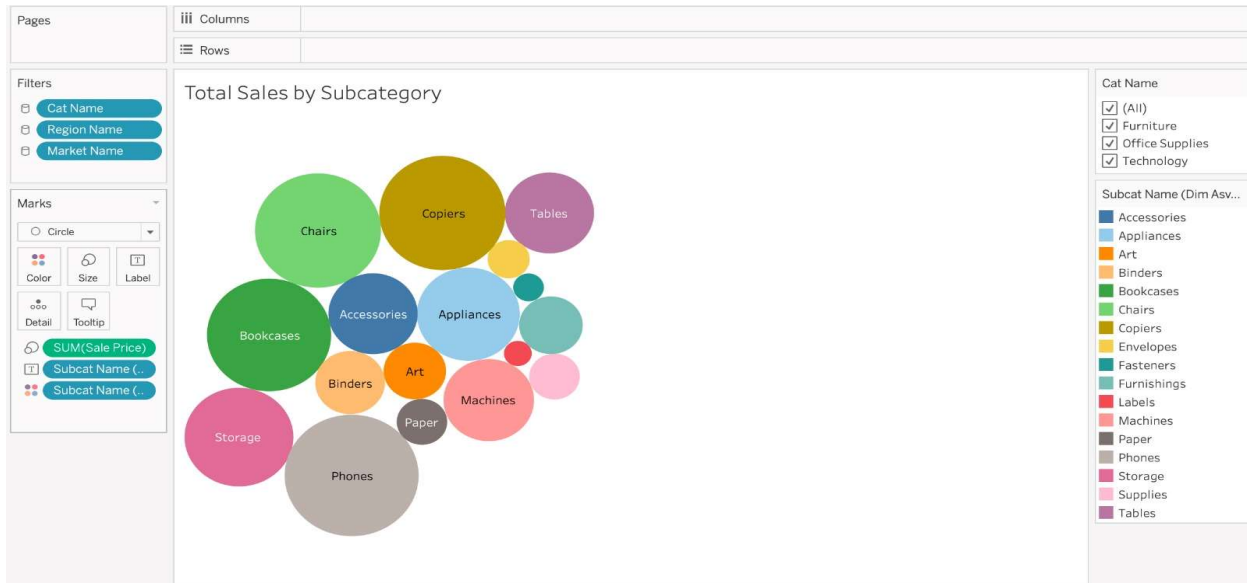
This type of graph is particularly useful because it allows us to quickly and easily identify which sub-categories are performing above or below average.





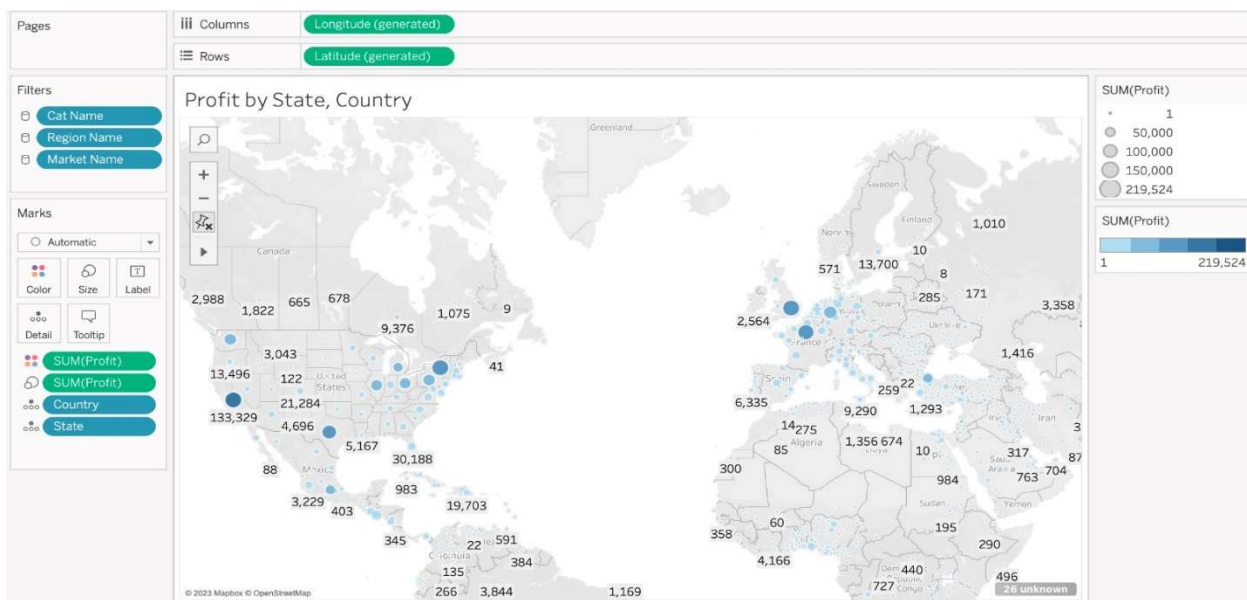
### Worksheet 3: Top sales by Sub-Category

Having a graph for this type makes it easier to quickly identify which sub-categories are driving the highest sales figures.



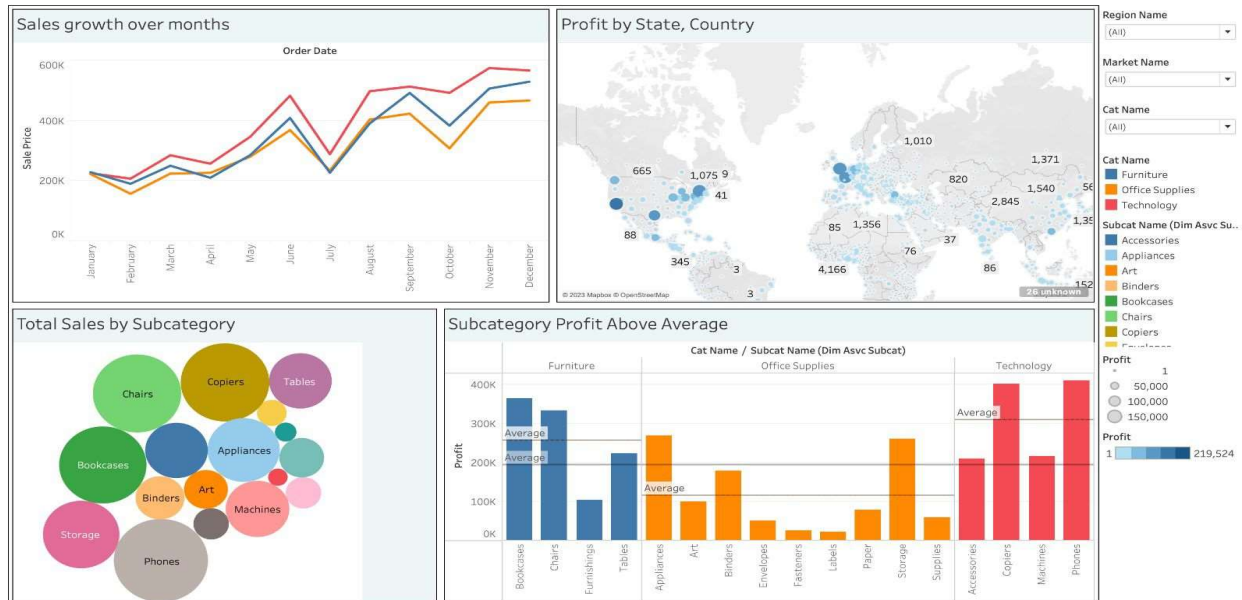
### Worksheet 4: Profit by States and Country.

This type of classification by states and country makes it easier to determine and analyze their profit performance in more detail. This can be especially useful for identifying areas of opportunity or concern, as well as for benchmarking performance against other regions or countries.



## Interactive dashboard:

We have created an interactive dashboard of multiple worksheets which will allow us to see how sales are changing over time for each region, market, category, etc. The graphs keep changing as we check or uncheck the filters on the right.



## PROJECT INSIGHTS

Throughout the duration of the project, we encountered many challenges and while solving the challenges, we almost always discovered and learned something new. We successfully learned how to use CSV files to load data in the OLTP database. It was challenging to map out the relations correctly and to join the tables in a way that represents the source data. Data Warehousing also presented its share of challenges and learnings, but we were able to learn how and why a DW is implemented. It was exciting to learn to represent data in the form of Fact and Dim tables so that required data could be viewed quickly with just a single join. Incremental ETL was also an intriguing concept as it allowed us to use the resources efficiently and only update the data which was updated or was not present in the DW before, which becomes really important as we scale our database as it is imperative that we do not transfer the same data again periodically. We also learned how to use the AWS S3 bucket to store data and then integrate it with the Snowflake data warehouse. One thing that the project made us realize is that it is not easy to make a flawless end-to-end data pipeline as we often found ourselves unraveling new information at a later stage which required modifications in the previous stages. Finally, after loading the data warehouse with the correct data, we learned how to use that data to provide insights to the end users and decision-makers. We grasped the concepts of Tableau to make relations within the DW tables after linking the platform with Snowflake. This allowed us to derive meaningful visualizations from the data which could be altered with filters according to the needs of the users. Overall, the experience was fruitful as it presented us with many challenges and learnings, and made us better at programming and collaboration in general.

## APPENDIX

### OLTP DDL:

```
CREATE TABLE asvc_address (cust_id
VARCHAR2(15) NOT NULL,
market_id VARCHAR2(4) NOT NULL,
street VARCHAR2(8) NOT NULL,
city VARCHAR2(15) NOT NULL,
state VARCHAR2(15) NOT NULL,
zipcode VARCHAR2(12) NOT NULL,
country VARCHAR2(20) NOT NULL,
segment VARCHAR2(20) NOT NULL,
tbl_last_dt DATE NOT NULL
);
```

```
COMMENT ON COLUMN asvc_address.street IS
'STREET OF THE ADDRESS OF CUSTOMER';
```

```
COMMENT ON COLUMN asvc_address.city IS
'CITY OF THE ADDRESS OF CUSTOMER';
```

```
COMMENT ON COLUMN asvc_address.state IS
'STATE OF THE ADDRESS OF CUSTOMER';
```

```
COMMENT ON COLUMN asvc_address.zipcode IS
'ZIPCODE OF THE ADDRESS OF CUSTOMER';
```

```
COMMENT ON COLUMN asvc_address.country IS
'COUNTRY OF THE ADDRESS OF CUSTOMER';
```

```
COMMENT ON COLUMN asvc_address.segment IS 'SEGMENT OF THE ADDRESS OF
CUSTOMER';
```

```
ALTER TABLE asvc_address ADD CONSTRAINT asvc_address_pk PRIMARY KEY (
market_id, cust_id);
```

```
CREATE TABLE asvc_ctgry (cat_id
VARCHAR2(4) NOT NULL, subcat_id
VARCHAR2(4) NOT NULL, cat_name
VARCHAR2(20) NOT NULL, descp
VARCHAR2(50), tbl_last_dt DATE
NOT NULL
);
```

```
COMMENT ON COLUMN asvc_ctgry.cat_id IS
'PRIMARY KEY FOR CATEGORY ENTITY';
```

```
COMMENT ON COLUMN asvc_ctgry.cat_name IS
'NAME OF THE CATEGORY';
```

```
COMMENT ON COLUMN asvc_ctgry.descp IS
'DESCRIPTION OF THE CATEGORY';
```

```
ALTER TABLE asvc_ctgry ADD CONSTRAINT asvc_ctgry_pk PRIMARY KEY (cat_id);
```

```
CREATE TABLE asvc_customer (cust_id
VARCHAR2(15) NOT NULL, cust_fname
VARCHAR2(10) NOT NULL,
cust_mname VARCHAR2(10), cust_lname
VARCHAR2(10) NOT NULL, tbl_last_dt
DATE NOT NULL
);
```

```
COMMENT ON COLUMN asvc_customer.cust_id IS
'PRIMARY KEY FOR CUSTOMER ENTITY';
```

```
COMMENT ON COLUMN asvc_customer.cust_fname IS
'FIRST NAME OF THE CUSTOMER';
```

```
COMMENT ON COLUMN asvc_customer.cust_mname IS
'CUSTOMER MIDDLE NAME';
```

```
COMMENT ON COLUMN asvc_customer.cust_lname IS
'CUSTOMERS LAST NAME';
```

```
ALTER TABLE asvc_customer ADD CONSTRAINT asvc_customer_pk PRIMARY KEY (
cust_id);
```

```
CREATE TABLE asvc_market (market_id
VARCHAR2(4) NOT NULL, region_id
VARCHAR2(4) NOT NULL, market_name
VARCHAR2(15) NOT NULL, tbl_last_dt
DATE NOT NULL
);
```

```
COMMENT ON COLUMN asvc_market.market_id IS
'PRIMARY KEY FOR MARKET ENTITY';
```

```
COMMENT ON COLUMN asvc_market.market_name IS
'MARKET NAME OF CUSTOMER ADDRESS LOCATED';
```

```
ALTER TABLE asvc_market ADD CONSTRAINT asvc_market_pk PRIMARY KEY (
market_id);
```

```

CREATE TABLE asvc_ordr (ord_r_id
VARCHAR2(15) NOT NULL, cust_id
VARCHAR2(15) NOT NULL, ord_r_date
DATE NOT NULL, ship_date DATE
NOT NULL, ord_r_priority VARCHAR2(15)
NOT NULL, modeof_ship
VARCHAR2(15) NOT NULL, tbl_last_dt
DATE NOT NULL
);

```

```

COMMENT ON COLUMN asvc_ordr.ord_r_id IS
'PRIMARY KEY FOR ORDER ENTITY';

```

```

COMMENT ON COLUMN asvc_ordr.ord_r_date IS
'DATE WHEN ORDER WAS PLACED';

```

```

COMMENT ON COLUMN asvc_ordr.ship_date IS
'SHIPPING DATE OR ORDER';

```

```

COMMENT ON COLUMN asvc_ordr.ord_r_priority IS
'PRIORITY OF ORDER';

```

```

COMMENT ON COLUMN asvc_ordr.modeof_ship IS
'MODE OF SHIPMENT';

```

```

ALTER TABLE asvc_ordr ADD CONSTRAINT asvc_ordr_pk PRIMARY KEY (ord_r_id);

```

```

CREATE TABLE asvc_ordr_prdct (
ord_r_id VARCHAR2(15) NOT NULL,

```

```
prdet_id VARCHAR2(15) NOT NULL,
qty NUMBER(4) NOT NULL,

 ship_cost NUMBER(4, 2) NOT NULL,
discount NUMBER(3, 2) NOT NULL,
profit NUMBER(4, 2) NOT NULL,
sale_price NUMBER(4, 2) NOT NULL,
tbl_last_dt DATE NOT NULL
);
```

```
COMMENT ON COLUMN asvc_ordr_prdet.qty IS
 'QUANTITY OF PRODUCTS';
```

```
COMMENT ON COLUMN asvc_ordr_prdet.ship_cost IS
 'SHIPPING COST ';
```

```
COMMENT ON COLUMN asvc_ordr_prdet.discount IS
 'DISCOUNT APPLIED ';
```

```
COMMENT ON COLUMN asvc_ordr_prdet.profit IS
 'PROFIT EARNED';
```

```
COMMENT ON COLUMN asvc_ordr_prdet.sale_price IS
 'SALES PRICE';
```

```
ALTER TABLE asvc_ordr_prdet ADD CONSTRAINT asvc_ordr_prdet_pk PRIMARY KEY (
ordr_id, prdet_id);
```

```
CREATE TABLE asvc_prdet (prdet_id
VARCHAR2(15) NOT NULL, cat_id
VARCHAR2(4) NOT NULL, prdet_name
```



```
VARCHAR2(15) NOT NULL,
prdt_descp VARCHAR2(50),
 tbl_last_dt DATE NOT NULL
);
```

```
COMMENT ON COLUMN asvc_prdt.prdt_id IS
 'PRIMARY KEY FOR PRODUCT
ENTITY';
```

```
COMMENT ON COLUMN asvc_prdt.prdt_name IS
 'NAME OF THE PRODUCT';
```

```
COMMENT ON COLUMN asvc_prdt.prdt_descp IS
 'DESCRIPTION OF THE PRODUCT';
```

```
ALTER TABLE asvc_prdt ADD CONSTRAINT asvc_prdt_pk PRIMARY KEY (prdt_id);
```

```
CREATE TABLE asvc_region (region_id
VARCHAR2(4) NOT NULL, region_name
VARCHAR2(15) NOT NULL, tbl_last_dt
DATE NOT NULL
);
```

```
COMMENT ON COLUMN asvc_region.region_id IS
 'PRIMARY KEY FOR REGION ENTITY';
```

```
COMMENT ON COLUMN asvc_region.region_name IS
 'REGION NAME WHERE MANY MARKETS ARE LOCATED';
```

```
ALTER TABLE asvc_region ADD CONSTRAINT asvc_region_pk PRIMARY KEY (region_id
);
```

```
CREATE TABLE asvc_return (
 rtn_id VARCHAR2(15) NOT NULL,
 ordr_id VARCHAR2(15) NOT NULL,
 returned CHAR(1) NOT NULL,
 tbl_last_dt DATE NOT NULL
);
```

```
COMMENT ON COLUMN asvc_return.rtn_id IS
 'PRIMARY KEY FOR RETURN ENTITY';
```

```
COMMENT ON COLUMN asvc_return.returned IS
 'ITEM IS RETURNED OR NOT';
```

```
ALTER TABLE asvc_return ADD CONSTRAINT asvc_return_pk PRIMARY KEY (rtn_id);
```

```
CREATE TABLE asvc_subcat (subcat_id
 VARCHAR2(4) NOT NULL, subcat_name
 VARCHAR2(15) NOT NULL, descp
 VARCHAR2(20), tbl_last_dt DATE NOT
 NULL
);
```

```
COMMENT ON COLUMN asvc_subcat.descp IS
 'DESCRIPTION OF SUBCATEGORY';
```

```
ALTER TABLE asvc_subcat ADD CONSTRAINT asvc_subcat_pk PRIMARY KEY (
 subcat_id
);
```

```
ALTER TABLE asvc_address
 ADD CONSTRAINT asvc_address_asvc_customer_fk FOREIGN KEY (cust_id)
 REFERENCES asvc_customer (cust_id);
```

```

ALTER TABLE asvc_address
 ADD CONSTRAINT asvc_address_asvc_market_fk FOREIGN KEY (market_id)
REFERENCES asvc_market (market_id);

ALTER TABLE asvc_ctgry
 ADD CONSTRAINT asvc_ctgry_asvc_subcat_fk FOREIGN KEY (subcat_id)
REFERENCES asvc_subcat (subcat_id);

ALTER TABLE asvc_market
 ADD CONSTRAINT asvc_market_asvc_region_fk FOREIGN KEY (region_id)
REFERENCES asvc_region (region_id);

ALTER TABLE asvc_ordr
 ADD CONSTRAINT asvc_ordr_asvc_customer_fk FOREIGN KEY (cust_id)
REFERENCES asvc_customer (cust_id);

ALTER TABLE asvc_ordr_prdct
 ADD CONSTRAINT asvc_ordr_prdct_asvc_ordr_fk FOREIGN KEY (ord_r_id)
REFERENCES asvc_ordr (ord_r_id);

ALTER TABLE asvc_ordr_prdct
 ADD CONSTRAINT asvc_ordr_prdct_asvc_prdct_fk FOREIGN KEY (prdct_id)
REFERENCES asvc_prdct (prdct_id);

ALTER TABLE asvc_prdct
 ADD CONSTRAINT asvc_prdct_asvc_ctgry_fk FOREIGN KEY (cat_id)
REFERENCES asvc_ctgry (cat_id);

ALTER TABLE asvc_return
 ADD CONSTRAINT asvc_return_asvc_ordr_fk FOREIGN KEY (ord_r_id)
REFERENCES asvc_ordr (ord_r_id);

```

## OLTP DML:

```
INSERT INTO asvc_ctgry (cat_id, cat_name,tbl_last_dt)
SELECT CONCAT('M', LPAD(ROW_NUMBER() OVER(ORDER BY Category), 3, '0')),
Category, NOW()
FROM (SELECT DISTINCT Category FROM orders_csv) AS t;
```

```
INSERT INTO asvc_market (market_id, market_name,tbl_last_dt)
SELECT CONCAT('MR', LPAD(ROW_NUMBER () OVER(ORDER BY Market), 3, '0')),
Market, NOW()
FROM (SELECT DISTINCT Market FROM orders_csv) AS t;
```

```
INSERT INTO asvc_subcat (subcat_id,cat_id, subcat_name,TBL_LAST_DT)
SELECT CONCAT('S', LPAD(ROW_NUMBER() OVER(ORDER BY subcat), 3, '0')),
m.cat_id, subcat,NOW()
FROM (SELECT DISTINCT subcat, Category FROM orders_Csv) AS t
INNER JOIN asvc_ctgry m ON m.cat_name = t.Category;
```

```
INSERT INTO asvc_region (region_id, market_id,region_name,tbl_last_dt)
SELECT CONCAT('R', LPAD(ROW_NUMBER() OVER(ORDER BY Region), 4,
'O')), m.market_id, Region, NOW()
FROM (SELECT DISTINCT Region, Market FROM orders_Csv) AS t
JOIN asvc_market AS m ON t.Market = m.market_name;
```

```
INSERT INTO asvc_returns (rtrn_id, ord_r_id, returned)
SELECT CONCAT('RE', LPAD(ROW_NUMBER() OVER (), 4, '0')), o.ord_r_id, r.returned
FROM returns_csv r
JOIN asvc_ordr o ON r.`Order ID` = o.ord_r_id;
```

```
INSERT IGNORE INTO asvc_ordr_prdct (ord_r_id, mode_of_ship, ship_date, QTY, ship_cost,
Discount, profit, sale_price, PRDCT_ID)
```

```

SELECT o.order_id, o.mode_of_ship, STR_TO_DATE(o.ship_date, '%m/%d/%Y'),
o.QUANTITY, CAST(REPLACE(o.ship_cost, '$', '') AS DECIMAL(10,5)), o.DISCOUNT,
o.Profit, CAST(REPLACE(o.Sales, '$', '') AS DECIMAL(10,5)), p.prduct_id
FROM orders_csv o
INNER JOIN ASVC_ORDR OD ON o.ORDER_ID = OD.ORDR_ID
INNER JOIN asvc_prduct p ON o.Product_ID = p.prduct_id;

```

```

INSERT INTO asvc_customer (cust_id, cust_fname, cust_lname, region_id,tbl_last_dt)
SELECT DISTINCT cust_id,
 SUBSTRING_INDEX(cust_name, ' ', 1) as cust_fname,
 SUBSTRING_INDEX(cust_name, ' ', -1) as cust_lname,
 r.region_id, NOW()
FROM orders_csv oc
JOIN asvc_region r ON oc.region = r.region_name;

```

```

INSERT ignore INTO asvc_address (cust_id, city, state, zipcode, country,
segment,tbl_last_dt,region_id)
SELECT c.cust_id, o.city, o.state, o.zipcode, o.country, o.segment, now(), r.region_id
FROM (
 SELECT DISTINCT CONCAT(cust_fname, ' ', cust_lname) AS cust_name, cust_id
 FROM asvc_customer
) c
JOIN orders_csv o ON o.cust_id = c.cust_id
JOIN asvc_region r ON o.region = r.region_name
WHERE NOT EXISTS (
 SELECT 1 FROM asvc_address a
 WHERE a.cust_id = c.cust_id
 AND a.city = o.city
 AND a.state = o.state
 AND a.zipcode = o.zipcode
 AND a.country = o.country
 AND a.segment = o.segment

```

```
 and a.region_id= r.region_id
);
```

```
INSERT ignore INTO asvc_ordr (ordr_id, cust_id, row_id, ordr_date, ordr_priority, tbl_last_dt)
SELECT DISTINCT o.order_id, c.cust_id, o.row_id,
STR_TO_DATE(o.order_date,'%m/%d/%Y'),
o.ordr_priority, now()
FROM orders_csv o
JOIN asvc_customer c ON o.cust_id = c.cust_id;
```

```
INSERT IGNORE INTO asvc_prdct (prdct_id, prdct_name, cat_id,tbl_last_dt)
SELECT p.p_id, p.prdct_name, c.cat_id,now()
FROM orders_csv p
JOIN asvc_ctgry c ON p.category =c.cat_name;
```

## OLTP DATA DICTIONARY QUERY

LIST OF TABLES:

```
1 • SELECT table_name
2 FROM information_schema.tables
3 WHERE table_schema = 'asvc';
4
5
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

| TABLE_NAME      |
|-----------------|
| asvc_address    |
| asvc_ctgry      |
| asvc_customer   |
| asvc_market     |
| asvc_ordr       |
| asvc_ordr_prdct |
| asvc_prdct      |
| asvc_region     |
| asvc_returns    |
| asvc_subcat     |
| orders_csv      |
| returns_csv     |
| ship_csv        |

LIST OF CONSTRAINTS FOR ALL TABLES:

```
1 • SELECT constraint_name, table_name, column_name, referenced_table_name, referenced_column_name
2 FROM information_schema.key_column_usage
3 WHERE referenced_table_schema = 'asvc';
4
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

| CONSTRAINT_NAME               | TABLE_NAME      | COLUMN_NAME | REFERENCED_TABLE_NAME | REFERENCED_COLUMN_NAME |
|-------------------------------|-----------------|-------------|-----------------------|------------------------|
| asvc_customer_asvc_region_fk  | asvc_customer   | region_id   | asvc_region           | region_id              |
| asvc_region_asvc_market_fk    | asvc_region     | market_id   | asvc_market           | market_id              |
| asvc_subcat_asvc_ctgry_fk     | asvc_subcat     | cat_id      | asvc_ctgry            | cat_id                 |
| asvc_address_asvc_customer_fk | asvc_address    | cust_id     | asvc_customer         | cust_id                |
| asvc_prdct_asvc_ctgry_fk      | asvc_prdct      | cat_id      | asvc_ctgry            | cat_id                 |
| asvc_prdct_asvc_subcat_fk     | asvc_prdct      | subcat_id   | asvc_subcat           | subcat_id              |
| asvc_ordr_asvc_address_fk     | asvc_ordr       | address_id  | asvc_address          | address_id             |
| asvc_ordr_asvc_customer_fk    | asvc_ordr       | cust_id     | asvc_customer         | cust_id                |
| asvc_ordr_prdct_asvc_ordr_fk  | asvc_ordr_prdct | ORDR_ID     | asvc_ordr             | ordr_id                |
| asvc_ordr_prdct_asvc_prdct_fk | asvc_ordr_prdct | prdct_id    | asvc_prdct            | prdct_id               |

## LIST OF COLUMNS AND COMMENTS FOR EACH TABLE:

```

1 • SELECT table_name, column_name, column_comment
2 FROM information_schema.columns
3 WHERE table_schema = 'asvc';
4

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

| TABLE_NAME    | COLUMN_NAME | COLUMN_COMMENT                     |
|---------------|-------------|------------------------------------|
| asvc_address  | address_id  |                                    |
| asvc_address  | street      | STREET OF THE ADDRESS OF CUSTOMER  |
| asvc_address  | city        | CITY OF THE ADDRESS OF CUSTOMER    |
| asvc_address  | state       | STATE OF THE ADDRESS OF CUSTOMER   |
| asvc_address  | zipcode     | ZIPCODE OF THE ADDRESS OF CUSTOMER |
| asvc_address  | country     | COUNTRY OF THE ADDRESS OF CUSTOMER |
| asvc_address  | segment     | SEGMENT OF THE ADDRESS OF CUSTOMER |
| asvc_address  | tbl_last_dt |                                    |
| asvc_address  | market_id   |                                    |
| asvc_address  | cust_id     |                                    |
| asvc_ctgry    | cat_id      | PRIMARY KEY FOR CATEGORY ENTITY    |
| asvc_ctgry    | cat_name    | NAME OF THE CATEGORY               |
| asvc_ctgry    | descp       | DESCRIPTION OF THE CATEGORY        |
| asvc_ctgry    | tbl_last_dt |                                    |
| asvc_ctgry    | subcat_id   |                                    |
| asvc_customer | cust_id     | PRIMARY KEY FOR CUSTOMER ENTITY    |
| asvc_customer | cust_fname  | FIRST NAME OF THE CUSTOMER         |

columns 11 x | Read

Output :



## OLAP DDL:

```
CREATE TABLE dim_asvc_ctgry (
 cat_id NUMBER(10) NOT NULL,
 cat_name VARCHAR2(20) NOT NULL,
 descp VARCHAR2(50)
);
```

```
ALTER TABLE dim_asvc_ctgry ADD CONSTRAINT dim_asvc_ctgry_pk PRIMARY KEY (
cat_id);
```

```
CREATE TABLE dim_asvc_customer (
 cust_id NUMBER(10) NOT NULL,
 address_id NUMBER(15) NOT NULL,
 cust_fname VARCHAR2(200) NOT NULL,
 cust_mname VARCHAR2(100),
 cust_lname VARCHAR2(200) NOT NULL,
 street VARCHAR2(200) NOT NULL,
 city VARCHAR2(100) NOT NULL,
 state VARCHAR2(100) NOT NULL,
 zipcode VARCHAR2(10) NOT NULL,
 country VARCHAR2(100) NOT NULL,
 segment VARCHAR2(50) NOT NULL
);
```

```
ALTER TABLE dim_asvc_customer ADD CONSTRAINT dim_asvc_customer_pk PRIMARY
KEY (address_id,

cust_id);
```

```
CREATE TABLE dim_asvc_date (
 date_id NUMBER(10) NOT NULL,
```

```
entry_id NUMBER(10),
full_date DATE,
days NUMBER(5),
month_short VARCHAR2(5),
month_num NUMBER(3),
month_long VARCHAR2(10),
day_of_week_short VARCHAR2(5),
day_of_week_long VARCHAR2(10),
year VARCHAR2(5),
quarter VARCHAR2(10)
);
```

```
ALTER TABLE dim_asvc_date ADD CONSTRAINT dim_asvc_date_pk PRIMARY KEY (
date_id);
```

```
CREATE TABLE dim_asvc_market (
market_id NUMBER(10) NOT NULL,
market_name VARCHAR2(15) NOT NULL
);
```

```
ALTER TABLE dim_asvc_market ADD CONSTRAINT dim_asvc_market_pk PRIMARY
KEY (market_id);
```

```
CREATE TABLE dim_asvc_prdct (
prdct_id NUMBER(10) NOT NULL,
prdct_name VARCHAR2(15) NOT NULL,
prdct_descp VARCHAR2(50)
);
```

```
ALTER TABLE dim_asvc_prdct ADD CONSTRAINT dim_asvc_prdct_pk PRIMARY KEY (
prdct_id);
```

```
CREATE TABLE dim_asvc_region (
 region_id NUMBER(10) NOT NULL,
 region_name VARCHAR2(20) NOT NULL,
 tbl_last_dt DATE NOT NULL,
 market_id VARCHAR2(4) NOT NULL
);
```

```
ALTER TABLE dim_asvc_region ADD CONSTRAINT dim_asvc_region_pk PRIMARY KEY
(region_id);
```

```
CREATE TABLE dim_asvc_subcat (
 subcat_id NUMBER(10) NOT NULL,
 subcat_name VARCHAR2(15) NOT NULL,
 descp VARCHAR2(20)
);
```

```
ALTER TABLE dim_asvc_subcat ADD CONSTRAINT dim_asvc_subcat_pk PRIMARY KEY
(subcat_id);
```

```
CREATE TABLE fact_asvc_order (
 "ROWNUM" NUMBER(15) NOT NULL,
 order_date DATE NOT NULL,
 order_priority VARCHAR2(20) NOT NULL,
 qty NUMBER(10) NOT NULL,
 ship_cost NUMBER(15) NOT NULL,
 ship_date DATE NOT NULL,
 discount NUMBER(3, 2) NOT NULL,
 mode_of_ship VARCHAR2(20) NOT NULL,
```

```

profit NUMBER(4, 2) NOT NULL,
sale_price NUMBER(4, 2) NOT NULL,
returned CHAR(1) NOT NULL,
date_id NUMBER(10) NOT NULL,
market_id NUMBER(10) NOT NULL,
region_id NUMBER(10) NOT NULL,
address_id NUMBER(15) NOT NULL,
subcat_id NUMBER(10) NOT NULL,
prdct_id NUMBER(10) NOT NULL,
cat_id NUMBER(10) NOT NULL,
cust_id NUMBER(10) NOT NULL
);

```

```

ALTER TABLE fact_asvc_order ADD CONSTRAINT fact_asvc_order_pk PRIMARY KEY (
"ROWNUM");

```

```

ALTER TABLE fact_asvc_order
ADD CONSTRAINT fact_order_dim_ctgry_fk FOREIGN KEY (cat_id)
REFERENCES dim_asvc_ctgry (cat_id);

```

```

ALTER TABLE fact_asvc_order
ADD CONSTRAINT fact_order_dim_customer_fk FOREIGN KEY (address_id,
 cust_id)
REFERENCES dim_asvc_customer (address_id,
 cust_id);

```

```

ALTER TABLE fact_asvc_order
ADD CONSTRAINT fact_order_dim_date_fk FOREIGN KEY (date_id)
REFERENCES dim_asvc_date (date_id);

```

```

ALTER TABLE fact_asvc_order
ADD CONSTRAINT fact_order_dim_market_fk FOREIGN KEY (market_id)

```

```
REFERENCES dim_asvc_market (market_id);
```

```
ALTER TABLE fact_asvc_order
```

```
ADD CONSTRAINT fact_order_dim_prdct_fk FOREIGN KEY (prdct_id)
REFERENCES dim_asvc_prdct (prdct_id);
```

```
ALTER TABLE fact_asvc_order
```

```
ADD CONSTRAINT fact_order_dim_region_fk FOREIGN KEY (region_id)
REFERENCES dim_asvc_region (region_id);
```

```
ALTER TABLE fact_asvc_order
```

```
ADD CONSTRAINT fact_order_dim_subcat_fk FOREIGN KEY (subcat_id)
REFERENCES dim_asvc_subcat (subcat_id);
```

## PROCEDURE CODE FOR INCREMENTAL ETL:

```
CREATE OR REPLACE PROCEDURE load_merge_fact_asvc_order()
RETURNS VARCHAR
LANGUAGE JAVASCRIPT
AS $$
 var err_code;
 var err_msg;
 try {
 var stmt = snowflake.createStatement({sqlText:`
 MERGE INTO fact_asvc_order a
 USING stg_incrmntl_fact_table b
 ON (a.ordr_id = b.ordr_id) AND (a.pr dct_id = b.pr dct_id)
 WHEN MATCHED THEN
 UPDATE SET
 a.order_date = b.order_date,
 a.order_priority = b.order_priority,
 a.qty = b.qty,
 a.ship_date = b.ship_date,
 a.ship_cost = b.ship_cost,
 a.mode_of_ship = b.mode_of_ship,
 a.discount = b.discount,
 a.profit = b.profit,
 a.sale_price = b.sale_price,
 a.address_id = b.address_id,
 a.market_id = b.market_id,
 a.region_id = b.region_id,
 a.cat_id = b.cat_id,
 a.cust_id = b.cust_id,
 a.subcat_id = b.subcat_id,
 a.subcat_name = b.subcat_name,
```

```

 a.returned = b.returned
 WHEN NOT MATCHED THEN
 INSERT (ordr_id, order_date, order_priority, qty, ship_date, ship_cost, mode_of_ship,
discount, profit,
 sale_price, address_id, market_id, region_id, prdct_id, cat_id, cust_id, subcat_id,
subcat_name, returned)
 VALUES (b.ordr_id, b.order_date, b.order_priority, b.qty, b.ship_date, b.ship_cost,
b.mode_of_ship, b.discount,
 b.profit, b.sale_price, b.address_id, b.market_id, b.region_id, b.prdct_id, b.cat_id,
b.cust_id, b.subcat_id, b.subcat_name, b.returned);
 `});
 var result = stmt.execute();
 return result.next() ? result.getColumnValue(1).toString() : 'Success';
} catch (err) {
 err_code = err.code;
 err_msg = err.message;
 return 'Error code ' + err_code + ': ' + err_msg;
}
$$;

CALL load_merge_fact_asvc_order();

```

## HISTORY TABLE TRIGGER

```
ALTER TABLE asvc_address_history ADD CONSTRAINT pk_asvc_address_history
PRIMARY KEY (address_id);
DELIMITER $$
CREATE TRIGGER tr_asvc_address_history BEFORE DELETE ON asvc_address
FOR EACH ROW
BEGIN
 INSERT INTO asvc_address_history SELECT * FROM asvc_address WHERE
address_id=OLD.address_id;
 UPDATE asvc_address_history
 SET tbl_last_dt=NOW()
 WHERE address_id=OLD.address_id;
END $$
DELIMITER ;
```

```
DELIMITER $$
CREATE TRIGGER tr_asvc_address_history3 after delete ON asvc_address
FOR EACH ROW
BEGIN
 INSERT INTO asvc_address_history SELECT * FROM asvc_address WHERE
address_id=OLD.address_id;
 UPDATE asvc_address_history
 SET tbl_last_dt=NOW()
 WHERE address_id=OLD.address_id;
END $$
DELIMITER ;
```

## TABLE PARTITIONING( CLUSTERING IN SNOWFLAKE)

```
CREATE TABLE partitioning_fact_asvc cluster by (MARKET_ID) AS
(SELECT * FROM fact_asvc_order);

SELECT SYSTEM$CLUSTERING_INFORMATION('partitioning_fact_asvc');

SHOW TABLES LIKE 'partitioning_fact_asvc';
```



## ETL CODE for dim\_asvc\_ctgry:

```
create schema asvc_project
```

```
CREATE OR REPLACE STORAGE integration aws_s3_integration
type = external_stage
storage_provider='S3'
enabled=true
storage_aws_role_arn= 'arn:aws:iam::649069303992:role/asvc-project'
storage_allowed_locations= ('s3://asvc-project/');
```

```
GRANT USAGE ON INTEGRATION aws_s3_integration TO ROLE accountadmin;
```

```
create or replace file format demo_format
type='CSV'
field_delimiter='|'
skip_header=1;
```

```
CREATE OR REPLACE STAGE demo_aws_stage
storage_integration = aws_s3_integration
file_format = demo_format
url='s3://asvc-project/';
```

```
CREATE TABLE dim_asvc_ctgry (
 cat_id VARCHAR2(10) NOT NULL,
 cat_name VARCHAR2(20) NOT NULL,
 descp VARCHAR2(50)
);
```

```
COPY INTO dim_asvc_ctgry
FROM @demo_aws_stage/PROJECT/asvc_ctgry.csv
FILE_FORMAT=(format_name=demo_format1 error_on_column_count_mismatch=false)
ON_ERROR = 'CONTINUE';
```